

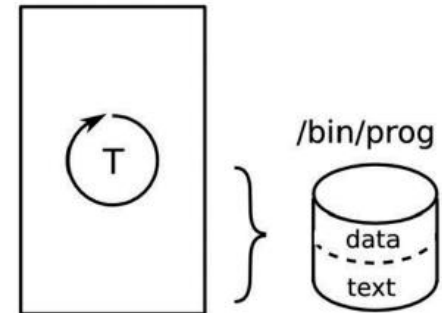
# EMBEDDED OPERATING SYSTEMS

---

Embedded Linux on Beaglebone Black

# Process

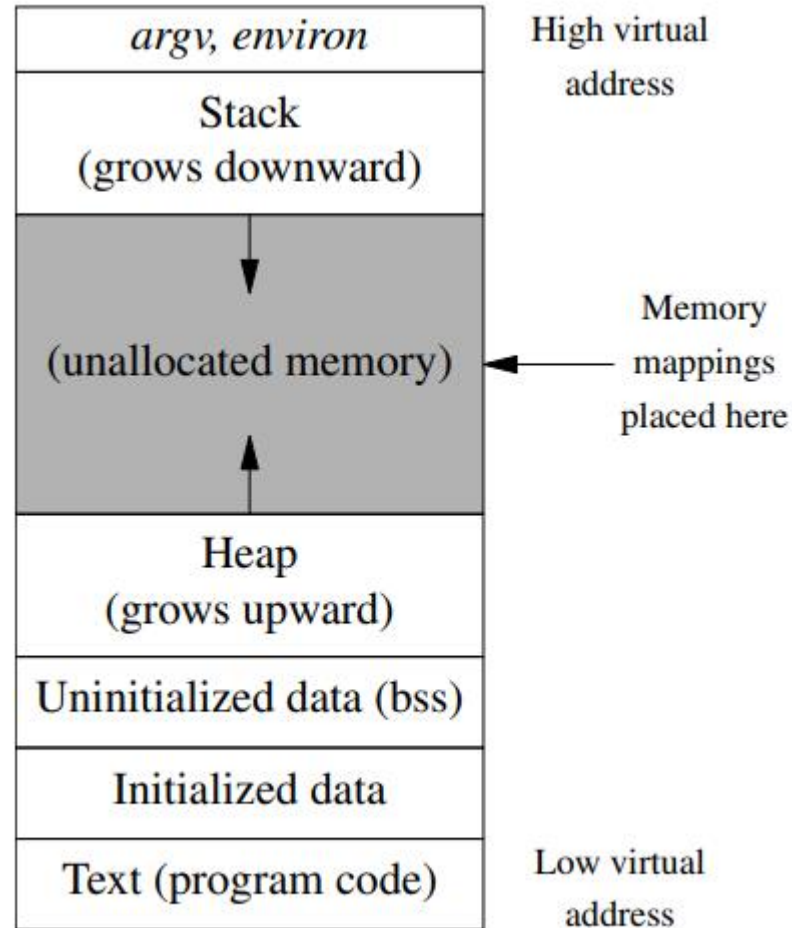
- A process is a combination of
  - Memory address space
  - Thread of execution
- Address space is private to the process
  - Other processes / threads cannot access it
  - Created by memory management subsystem in kernel
- Process has resources
  - Memory (stack, heap), file descriptors, etc.
- Kernel reclaims all resources when process ends
- Processes communicate among themselves using IPC
  - Inter-process communication



# Process memory space

- Divided into segments
  - **Text:**
    - Machine-language instructions
    - Marked read-only to prevent modifications
    - Multiple processes could share same code
  - **Initialized data**
    - Global and static variables explicitly initialized
    - Valued read from program when it is loaded
  - **Uninitialized data**
    - Global and static variables that are not explicitly initialized
    - Initialized to zero when process is created
  - **Stack**
    - Storage for function locals and linkage info
  - **Heap**
    - Area from which memory can be dynamically (de)allocated

# Process memory layout



# Process environment

- Each process has a list of environment variables
  - Strings of the form name=value
- New process (child) inherits parent's environ
  - Simple 1-way IPC!
- Commonly used to control program behavior
- Examples:
  - USER: User's name
  - HOME: User's home directory
  - SHELL: User's shell
  - PATH: User's path

# The /proc file system

- Psuedo-file system
  - Exposes kernel information
  - Via file system metaphor
  - Structured as set of subdirectories and files
- Files don't really exist
  - Created on the fly in SDRAM
  - Many files are read-only
  - Some files are writable
    - Can update / change kernel settings

# Examples: /proc file system

- `/proc/cmdline`: command line used to start kernel
- `/proc/cpuinfo`: info about CPUs on the system
- `/proc/meminfo`: info about memory and memory usage
- `/proc/modules`: info about loaded kernel modules
- `/proc/sys/fs/`: files and subdirectories with filesystem-related info
- `/proc/sys/kernel/`: files and subdirectories with various readable/settable kernel parameters
- `/proc/sys/net/`: files and subdirectories with various readable/settable networking parameters

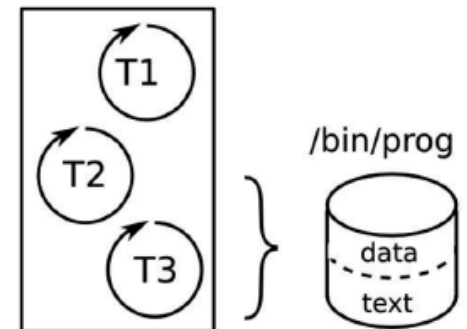
# /proc/PID directories

- Each running process has a /proc/PID directory
  - Contains subdirectories and files
    - Exposing information about process with that PID
    - Examples:
      - cmdline: command line used to start the process
      - cwd: current working directory
      - environ: environment of current process
      - fd: info on open file descriptors
      - limits: resource limits
      - mounts: devices mounted by that process
      - status: info about the process



# Thread

- A thread is an executional entity within a process
- All processes begin with *main()* thread
  - They may create more threads using **pthread** APIs
- Multiple threads in same process share
  - Same memory address space
  - Other resources like file descriptors
- Communication is easier
  - Needs to be properly synchronized



# Creating a process: fork()

- Processes are created using the **fork()** call

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(void);
```

- Here pid\_t is the Process ID

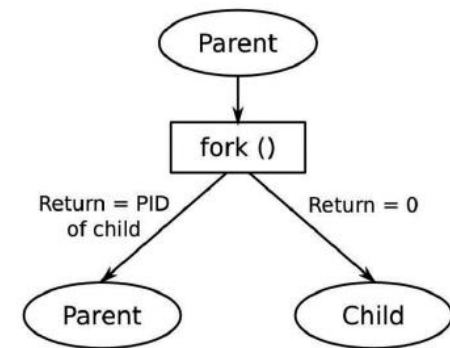
- But fork() is a strange function

- 2 returns on success

- First return is 0 in the new process created (child)
- Second return is a number in the original process (parent)

- Child is an exact copy of parent

- But in separate address spaces!



# Terminating a process: exit()

- Process terminates when
  - It calls exit() function – **voluntary**
  - It gets a signal (like SIGKILL) – *involuntary (is killed)*
- Return value of terminating process is either
  - Argument to exit() call
  - Signal value if killed
- Return value can be collected by parent
  - Using **wait()** or **waitpid()** calls
  - Kernel can notify parent using a SIGCHLD
- If parent is killed / not waiting
  - Child process becomes a **zombie**
  - Attached to init / kernel

# Forking exercise

- ***forking.c***

- Notice how the code is written
  - Common codebase for child and parent
  - How child / parent process can be distinguished
    - Trapping the return value of the fork() call
  - How child process exits with an exit code
  - How parent can wait for child termination
    - And get access to exit code returned by child

# Forking different program: `exec()`

- What if we want a `fork()` but want to run something else, instead of parent code?
  - Use the `exec()` family of functions
    - Get ***location*** of the executable as argument
      - `execl()`
      - `execvp()`
      - `execle()`
    - Get a ***vector*** of arguments
      - `execv()`
      - `execvp()`
      - `execve()`

# The exec() family

- *int execl(const char \*pathname, const char \*arg, ..., NULL);*
- *int execlp(const char \*file, const char \*arg, ..., NULL);*
- *int execl\_e(const char \*pathname, const char \*arg, ..., NULL, char \*const envp[]);*
- *int execv(const char \*pathname, char \*const argv[]);*
- *int execvp(const char \*file, char \*const argv[]);*
- *int execve(const char \*pathname, char \*const argv[], char \*const envp[]);*

# exec() exercise

- Use the shell script (run\_all\_execs.sh)
  - To run all types of exec() calls
- Note the different calls
  - Input parameters and usage!
- Terminology
  - *l* → *location*
  - *v* → *vector*
  - *p* → *path*
  - *e* → *environment (shell)*

# Inter-Process Communication

- Each process is an island of memory
- Inter-Process Communication (IPC) happens
  - By copying info from one address space to another
    - Using a queue / buffer to hold and pass messages
      - Examples: **sockets**, **pipes**, **message queues**
    - Using a **shared file** and **file locks**
  - By creating a memory area that all can access
    - Using shared memory; needs access synchronization
      - Examples: **shmem**
- Message and shared file based methods are easier to program and debug
  - But slow for large number of messages
- Shared memory based methods are more efficient
  - But need careful synchronization



# IPC using shared files, file locks

- Writer and reader process
  - Share a common file on file system (file.dat)
- Either operation needs a file lock (***flock***)

```
struct flock {  
    ...  
    short l_type;    /* Type of lock: F_RDLCK,  
                     F_WRLCK, F_UNLCK */  
    short l_whence;  /* How to interpret l_start:  
                     SEEK_SET, SEEK_CUR, SEEK_END */  
    off_t l_start;   /* Starting offset for lock */  
    off_t l_len;     /* Number of bytes to lock */  
    pid_t l_pid;     /* PID of process blocking our lock  
                     (set by F_GETLK and F_OFD_GETLK) */  
    ...  
};
```

- Lock types:
  - Writer needs a F\_WRLCK (exclusive)
  - Reader needs a F\_RDLCK (shared / non-exclusive)
  - After the operation, process needs to F\_UNLCK

# File locking using `fcntl()`

- Locks are obtained using **`fcntl()`**

```
#include <unistd.h>
```

```
#include <fcntl.h>
```

```
int fcntl(int fd, int cmd, ... /* arg */);
```

- Commands for `fcntl` (*cmd*):

- GETLK

- Get the current lock; returns immediately

- SETLK

- Set the lock (3<sup>rd</sup> argument is `&lock`); no waiting

- SETLKW

- Set the lock (3<sup>rd</sup> argument is `&lock`), wait until you get the lock

# File lock exercise

- *flock-writer.c*
  - Implements the writer
  - Creates the file, and takes a write lock (F\_WRLCK) – with wait
  - Writes data to file
  - Exits after releasing the file lock (F\_UNLCK)
- *flock-reader.c*
  - Implements the reader
  - Takes a read lock (F\_RDLCK) on the file
    - Waits till it gets it!
  - Reads data from file and displays on stdout
  - Exits after releasing the lock (F\_UNLCK)
- Run *flock-writer.out* and *flock-reader.out* on different terminals at once
- Try running 2 writer processes (*flock-writer.out*) at once

# IPC using shared memory

- POSIX API for sharing memory
  - *shm\_open()*
    - Create/open a new/existing shared memory object
  - *ftruncate()*
    - Set the size of the shared memory object (viewed as fd)
  - *mmap()* / *munmap()*
    - Map/unmap devices or files into process memory space
  - *shm\_unlink()*
    - Unlink (delete) share memory object
- POSIX maintains a backing file for shared memory
  - Located in /dev/shm/...
- We need to use primitives like *semaphore/mutex*
  - For controlling exclusive access to memory object
- Linking flags: ***-lrt -lpthread***

# POSIX shared memory API

- Open/create and close

```
#include <sys/mman.h>
#include <sys/stat.h>      /* For mode constants */
#include <fcntl.h>         /* For O_* constants */
int shm_open(const char *name, int oflag, mode_t mode);
int shm_unlink(const char *name);
```

- Set size

```
#include <unistd.h>
#include <sys/types.h>
int ftruncate(int fd, off_t length);
```

- Memory map/unmap

```
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);
int munmap(void *addr, size_t length);
```

# POSIX semaphore API

- Open/create

```
#include <fcntl.h>          /* For O_* constants */
#include <sys/stat.h>       /* For mode constants */
#include <semaphore.h>
sem_t *sem_open(const char *name, int oflag, mode_t mode, unsigned int value);
```

- Semaphore mode
  - **O\_CREAT | O\_RDWR** for writer
  - **O\_RDWR** for reader

- Wait on / Post a semaphore

```
int sem_wait(sem_t *sem);
int sem_post(sem_t *sem);
```

- Close/unlink

```
int sem_close(sem_t *sem);
int sem_unlink(const char *name);
```

# Shared memory exercise

- *shmem-writer.c*
  - Implements the writer
  - Opens/creates and mmap's the shared mem segment
    - Sets the size
  - Creates the semaphore
    - Writes to mem segment while holding it
  - Releases (posts) the semaphore
  - Cleans up (semaphore deletion, unlinking backing file, etc.)
- *shmem-reader.c*
  - Implements the reader
  - Opens and mmap's the shared mem segment
  - Tries to take the semaphore
    - Waits till it gets the semaphore
  - Reads from the segment and outputs to stdout
  - Unmaps the mem segment and exits

# Threads on Linux

- Linux uses the POSIX threads library
  - Commonly called ***pthread***
- Threads are called 'light-weight processes' (LWP)
  - Share same resources (memory space, file descriptors) as the parent process
- A process starts by creating a single thread
  - Called main thread (from the C ***main()*** entryptoint)
- Inter-thread communication
  - Is simpler and has less overhead
  - But needs to be synchronized and controlled!
- Linking: ***-lpthread***



# What threads share (& don't!)

- Threads in a process share:
  - Process ID (PID)
  - Parent Process ID (PPID)
  - User ID (UID) and Group ID (GID)
  - Controlling terminal
  - File descriptors and locks
- Distinct for each thread:
  - Thread ID (TID)
  - Signal mask
  - ***errno*** variable
  - RT scheduling policy and priority

# POSIX thread API

- Open/create

*#include <pthread.h>*

*int **pthread\_create**(pthread\_t \*thread, const pthread\_attr\_t \*attr, void \*(\*start\_routine) (void \*), void \*arg);*

- Termination options

*void **pthread\_exit**(void \*retval);*

*int **pthread\_cancel**(pthread\_t thread);*

*int **pthread\_join**(pthread\_t thread, void \*\*retval);*

- Return from start\_routine()

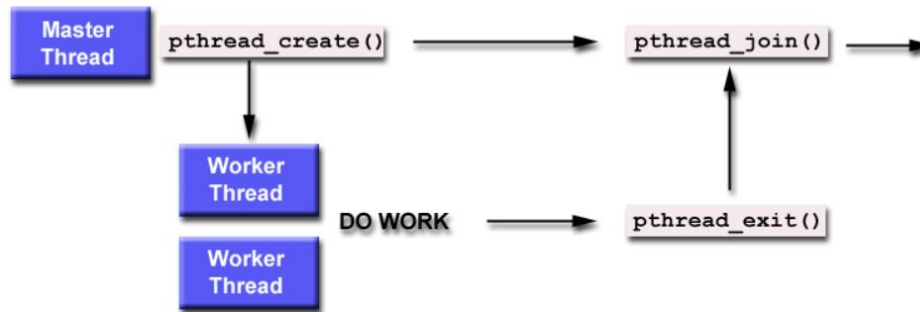
- Thread attributes

*int **pthread\_attr\_init**(pthread\_attr\_t \*attr);*

*int **pthread\_attr\_destroy**(pthread\_attr\_t \*attr);*

# Joinable vs. Detachable threads

- Usual lifecycle of threads



- By default, threads are joinable
  - **Master** waits for **worker** using `pthread_join()`
  - **Worker** resources not freed up till **master** joins, even if **worker** exits
- But what if we want to free **worker** resources early on?
  - Then **worker** can be created **detachable**

# Creating detachable threads

- Create a ***pthread\_attr\_t attr***;
- Set/get detach state using these APIs:  
*int pthread\_attr\_setdetachstate(pthread\_attr\_t \*attr, int detachstate);*  
*int pthread\_attr\_getdetachstate(const pthread\_attr\_t \*attr, int \*detachstate);*
- Detach states:
  - *PTHREAD\_CREATE\_DETACHED*
  - *PTHREAD\_CREATE\_JOINABLE*
- Create the thread using this ***attr***  
*int pthread\_create(pthread\_t \*thread, const pthread\_attr\_t \*attr, void \*(\*start\_routine) (void \*), void \*arg);*
- A joinable thread can be changed to detached using the API  
*int pthread\_detach(pthread\_t thread);*

# Joinable/detached exercise

- Run *join-detach.out*
- Observe the output
  - When threads are created
  - As threads execute
  - When main tries to join the threads
- Notice that
  - Joinable thread can be joined
  - Detached threads cannot be joined any more

# Mutex

- Mutex is an exclusive lock
  - Used for controlling / serializing access
  - To a common (shared) resource
  - Between multiple threads
- Only the thread holding the lock
  - Can access the shared resource
  - It then releases (unlocks) the mutex
- The other thread has to wait till the mutex is free

# Conditional variable (condvar)

- Threads need to be able to alert other threads
  - Something has changed / needs attention
  - This is called a condition
  - The alert is sent through a ***conditional variable (condvar)***
- One thread (alerter) signals the other (alertee)
  - By signaling using the ***condvar***
- The other thread sleep-waits on the ***condvar***
  - Waiting for the signal from the 1<sup>st</sup> thread
- Since both threads use/access the ***condvar*** at one time
  - We need a ***mutex*** to serialize access to the ***condvar***

# POSIX mutex and condvar API

- Mutex API

- Creation

- pthread\_mutex\_t* mutex = PTHREAD\_MUTEX\_INITIALIZER;

- Lock/unlock

- int pthread\_mutex\_lock(pthread\_mutex\_t \*mutex);*

- int pthread\_mutex\_unlock(pthread\_mutex\_t \*mutex);*

- Conditional variable (condvar) API

- Creation

- pthread\_cond\_t* cond = PTHREAD\_COND\_INITIALIZER;

- Wait/signal

- int pthread\_cond\_wait(pthread\_cond\_t \*cond, pthread\_mutex\_t \*mutex);*

- int pthread\_cond\_signal(pthread\_cond\_t \*cond);*



# Mutex-Condvar exercise

- Run *mutex-condvar.out*
- Observe
  - Producer and Consumer share a global char buffer
  - Producer enters data into the buffer
    - Signals this to Consumer
  - Consumer waits for signal from Producer
    - Retrieves data from buffer; resets it to “empty”
  - They use a *condvar* and a *mutex*
    - For serialized and controlled access to the *condvar*

# THANK YOU!

---