

EMBEDDED OPERATING SYSTEMS

Embedded Linux on Beaglebone Black

Scheduling

- Scheduling
 - Process of assigning **resources** to perform **tasks**
- Resources
 - CPU, memory, network links, registers, buses
- Tasks
 - Processes, threads
- Carried out by a **Scheduler**
- Schedulers aim to:
 - Achieve optimum resource utilization
 - Allow resource sharing between multiple users
 - Achieve a target quality of service

Scheduler Goals

- A Scheduler may aim at one/more goals:
 - Maximizing throughput (work done per unit time)
 - Minimizing wait times (time for work to get started)
 - Minimizing latency (system response time)
 - Maximizing fairness (distributing CPU in a just manner)
- In **real-time** environments
 - Schedulers have to meet ***deadlines***
 - Time by which an activity **should** be complete

Scheduling in Linux

- The Linux scheduler
 - Has a queue of threads all ready to run
 - It schedules them on CPU(s) as they become available
- Scheduling is on ***thread basis***
 - Not on process basis!
- Each thread has a ***scheduling policy***
 - Either ***real-time*** or ***time-shared***
- Real-time threads have ***priority***
 - Concept of pre-emption
- Time-shared threads have ***niceness***
 - Concept of entitlement to CPU time

When does scheduler run?

- The Linux scheduler runs when:
 - A thread is blocked by calling sleep()
 - Or any other blocking call
 - A time-shared thread exhausts its time-slice
 - An interruption causes a thread to be unblocked
- Linux seeks a balance between
 - Fairness (by time-sharing)
and
 - Determinism (by using real-time)

Time-sharing vs. Real-time

- Time-sharing
 - Based on principle of fairness
 - Each thread should get its fair amount of CPU time
 - No thread should be hog the system
 - If a thread runs for too long, it is put to the back of the queue
 - Threads doing more work are given more resources
 - Automatically adjust to varying workloads
- Real-time
 - Based on concept of priority
 - When the demand is for deterministic behavior
 - Minimal guarantee that thread will meet its deadline
 - Real-time threads pre-empt time-shared threads
 - Scheduler decides which to run based on priority
 - Most RTOS schedulers are built this way

Time-shared Policies

- CFS (Completely Fair Scheduler)
 - Scheduler used for time-sharing
 - Calculates a running tally of time consumed by a thread
 - Compares it with its **entitlement** and adjusts accordingly
 - A thread exceeding entitlement is suspended
- Policies
 - SCHED_NORMAL (SCHED_OTHER)
 - Default policy, most threads run with this policy
 - SCHED_BATCH
 - Time granularity is larger, to reduce context switching overheads
 - SCHED_IDLE
 - Run only when no other threads run; lowest possible priority

Time-sharing: Niceness

- Some time-shared threads are more important
 - Indicated by its ***niceness*** property
 - A number between 19 and -20
 - 19 is really nice, -20 is really not nice
 - A thread is nice if its loads the CPU less!
 - Nice value can be changed for
 - SCHED_NORMAL and SCHED_BATCH threads
 - Only root user can change the niceness property
 - Bash has commands: ***nice, renice***

Real-time Policies

- Meant for achieving determinism
 - Run the highest-priority real-time when it is ready
 - Real-time always pre-empts time-shared
 - Priority in the range of 1 (lowest) to 99 (highest)
- Policies
 - SCHED_FIFO
 - Run to completion
 - Once started, thread runs unless
 - Preempted / blocked / terminated
 - SCHED_RR
 - Round robin between threads at same priority
 - By time slicing (usually 100ms)

Balance: Real-time – Time-share

- Problem with real-time
 - A **rogue** real-time process with **high priority**
 - Can prevent
 - All lower priority real-time threads
 - All time-shared threads
 - From running / getting resources
- System ***locks up*** and becomes erratic
- Linux scheduler solution:
 - Reserve 5% of time
 - For non-real-time threads
 - Configurable using
 - `/proc/sys/kernel/sched_rt_period_us`
 - `/proc/sys/kernel/sched_rt_runtime_us`
 - Use watchdogs

Choice of Policy

- For most systems, time-share works fine
- A thread may qualify for real-time if:
 - It has a strict deadline for deliverable
 - Missing the deadline compromises system efficacy
 - It is event-driven
 - It is not compute-bound (CPU hog)
- Choice of RT policy
 - Rate Monotonic Analysis (RMA)

THANK YOU!
